


Malloc - Presentation

ACU 2020 Team



I never really sleep. Got one eye open, always.

This document is for internal use only at EPITA <<http://www.epita.fr>>.

Copyright © 2019-2020 Assistants <assistants@tickets.assistants.epita.fr>.

Rules


- You must have downloaded your copy from the Assistants' Intranet <<https://intra.assistants.epita.fr>>.
- This document is strictly personal and must **not** be passed on to someone else.
- Non-compliance with these rules can lead to severe sanctions.

Introduction

What's different from other algorithms ?

- Low fragmentation.
- Mitigated performances on its own.
- Highly depends on the complementary strategies used with this algorithm.
- Really challenging approach.

- The idea is to manage a contiguous memory block of size n with n a power of two.
- The block can be given to the user as is or divided into two sub-blocks.
- The two sub-blocks are called buddies (each one is the buddy of the other).
- The logic is recursive and each sub-block can be given to the user or divided as well into two sub-block, etc...
- Each block's size is therefore a power of two.
- You have to define a minimum size for the blocks.



Management

- User does a malloc of size s .
- Compute the first power of two greater than or equal to s : $s2p$.
- Of course $s2p$ is greater or equal to the chosen minimum size.
- From there, find the smallest block possible of size at least $s2p$ in the tree.
- If the free block found is at least 2 times greater than $s2p$, divide it into two sub-block (the division is done recursively) until the sub-block size is as small as possible and still fits $s2p$.

- Mark the block as free.
- If the buddy of the block is also available, merge them to create a larger block.
- The merge is done recursively until it is not possible anymore.


Start with an initial block of size 1024 bytes.



(free 1024)

Allocation of 512 bytes:

- Split the block in two blocks of 512 bytes and use the first one.



Step 3

Allocation of 64 bytes:

- Split the free block of 512 bytes into two blocks of 256 bytes.
- Split the first block of 256 bytes into two blocks of 128 bytes.
- Split the first block of 128 into two blocks of 64 bytes and use the first one.



Allocation of 120 bytes:

- The rounding power of two is 128 bytes
- Use the already available block of 128 bytes.

512	64	64	128	256
-----	----	----	-----	-----

512	64	64	128	256
-----	----	----	-----	-----


Allocation of 500 bytes:

- The rounding power of two is 512 bytes.
- There is no block large enough: allocation is not possible.
- Would have to create a new binary buddy (not in this example).



Free the 512 bytes' block:

- Mark the block as free.
- The buddy of the block is not free so no merge is needed.



Step 7

Free the 64 bytes' block:


- Mark the block as free.
- It's buddy is also free: merge them into a free block of 128 bytes.
- The buddy of the new 128 bytes block is not free, the merge is not possible.



Step 8



Free 128 bytes' block:

- Mark the block as free.
- Its buddy is free: merge them into a block of 256 bytes.
- The buddy of the new 256 bytes' block is free: merge them into a block of 512 bytes.
- The buddy of the new 512 bytes' block is free: merge them into a block of 1024 bytes.



At the beginning of each block you need some metadata, with at least:

- block status (allocated / free)
- The block size




Start by reading the metadata on the left (first block):

- If the block is not available, move the pointer from block's size bytes to check the next block.
- If the block is of the right size you just need to mark it as used.
- If the block is available but too big, it needs to be split:
 - Update its metadata (reduce block size).
 - Create new metadata for the buddy issued by the division.
- The buddy of a block can be easily accessed with the formula:
 - $\text{block_addr} \oplus \text{block_size}$

Note: With the metadata, a block of size n (n a power of two) can fit for an allocation of size $n - \text{sizeof}(\text{metadatas})$.

- Allocation is very slow, while freeing a block is fast.
- It's almost necessary to use a datastructure to keep track of the free blocks.
- Various datastructure might be used : sized free-lists, binary tree..

Example using sized free-lists



Conclusion

Any questions?